Natural Language Processing
Assignment 1

Efe Özbal (s5996821)
Sigurdur Haukur Birgisson (s5982960)
Josh van Vliet (s5950465)

February 20, 2026

1 Dataset & Setup

Dataset for this text classification by topic task is the textual AG News dataset from Hugging Face [Zhang et al.(2015)Zhang, Zhao, and LeCun]. Each data point has an associated label, the news title and its description. There are 4 classes, all of which has 30000 training and 1900 testing samples: world, sports, business, and sci/tech.

We use the official Hugging Face split (120,000 train and 7,600 test). We use stratified sampling (90/10) with seed 67 to create a development set:

- Train: 108,000
- Dev: 12,000
- Test: 7,600

1.1 Preprocessing

The input is the title and description concatenated. We used the following preprocessing steps:

- Lowercasing
- Whitespace removal
- Punctuation removal
- English stopword removal (using NLTK library)
- Lemmatization (with WordNetLemmatizer)

The reason why we remove stop words, is because common words often carry less information and carry noise into our classifier.

The last step in the normalization is lemmatization, the act of reducing words into their dictionary form (e.g. "running" to "run"). Using the NLTK's WordNetLemmatizer, we reduced each word in the text to their dictionary forms to reduce the computational complexity. In this way, different forms of the same word are reduced to a single form.

1.2 Tokenization

Even though there is no explicit tokenization in our code, texts are vectorized with word-level TF-IDF. After the normalization steps, TF-IDF vectorizes, and thus tokenizes, the input text into numerical data:

- n-gram range: (1, 2)
- `max_features` = 10,000
- `min_df` = 2

1.3 Reproducibility

All experiments make use of the fixed random seed 67. We use it for the data split and model initialization where it is applicable.

2 Models

In this section, we describe the two models we implemented for the text classification task:

- TF-IDF + Logistic Regression
- TF-IDF + LinearSVM.

Both models utilize Term Frequency-Inverse Document Frequency (TF-IDF) for feature extraction, which transforms the textual data into a numerical representation based on the importance of words in the documents.

The only difference between them is in the linear classifier head.

2.1 Hyperparameter Search

Model selection was done only on the dev split.

2.1.1 Logistic Regression Grid

- $C \in \{0.01, 0.1, 1.0, 10.0, 100.0\}$
- `max_iter` = 1,000

2.1.2 Linear SVM Grid

- $C \in \{0.01, 0.1, 1.0, 10.0, 100.0\}$
- loss $\in$ {hinge, squared_hinge}

2.1.3 Best Dev Split Settings

The best settings selected on the dev split were:

- Logistic Regression: $C = 1.0$ (Dev Accuracy: 0.9145)
- Linear SVM: $C = 1.0$, loss=hinge (Dev Accuracy: 0.9163)

The model architecture is shown in Figure 1. The input text is first preprocessed and then transformed into TF-IDF features. These features are then fed into either a Logistic Regression classifier or a Linear Support Vector Machine (SVM) for classification.

3 Experimental Protocol

In this section, we describe the experimental protocol we followed for tuning our models and selecting the best hyperparameters.

3.1 Model Selection Protocol

We followed this procedure for model selection:

1. Build vocabulary for TF-IDF on training split only.
2. Train each model on train/dev splits.
3. Select the best hyperparameters based on our dev split accuracy.
4. Do a final evaluation on the test split for final results.

This protocal ensures that we do not have any data leakage from our test data into our model selection.

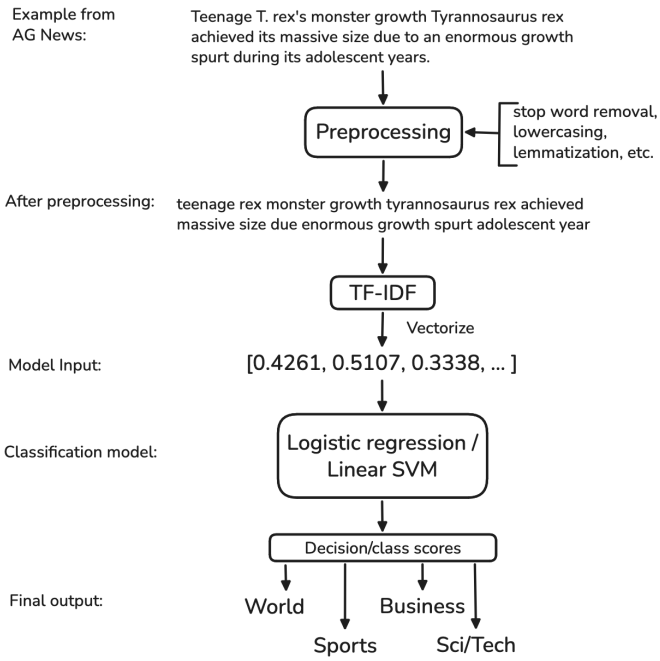

Figure 1: Model architecture for TF-IDF + Logistic Regression and TF-IDF + LinearSVM. The figure shows a random sample from the training set, which is preprocessed and transformed into TF-IDF features before being fed into the classifiers.

3.2 Training and Reproducibility

We set max_iter = 1,000 for Logistic Regression and Linear SVM as early stopping. This is to prevent unnecessary optimization steps and promote convergence.

During training, we kept the random seed at 67 (same as for data splitting) for reproducibility purposes.

3.3 Hyperparameter Tuning

Hyperparameters were tuned with grid search on just the dev split. The exact details are defined in Section 2.

3.4 Evaluation and Fair Comparison

For each model, we look at:

- Accuracy

- Macro-F1
- Confusion matrix

Macro-F1 is used to show performance for all four classes.

To make sure there is a fair comparison between the models, both have:

- The same train/dev/test splits
- The same preprocessing and TF-IDF feature settings
- The same random seed (67)
- Tuning only on the dev split and a final test eval

4 Results

In this section, we present the results of our experiments on the test set for both the Linear SVM and Logistic Regression models. We report the classification reports (precision, recall, F1-score, and support) for each class, as well as the confusion matrices for both models. Additionally, we provide a comparison of the overall accuracy and macro-F1 scores for both models on the test set.

We hyperparameter tuned both models on the dev set, selecting the best parameters based on the highest accuracy. For the Linear SVM, the best parameters were C=1.0 and loss='hinge', while for Logistic Regression, the best parameter was C=1.0. We then evaluated both models on the test set using these best parameters and report the results below.

In Table 1, we can see that the Linear SVM model achieved an overall accuracy of 0.92 and a macro-F1 score of 0.92 on the test set. The Logistic Regression model, as shown in the subsequent table (Table 2), achieved an overall accuracy of 0.91 and a macro-F1 score of 0.91 on the test set.

In the confusion matrices (Figures 2 and 3), we can observe that both models have similar patterns of misclassification, with the most common confusion occurring between the sci/tech and business classes. This is reflected in the classification reports where the precision, recall, and F1-scores for these classes are lower compared to the sport and world classes. Overall, the Linear SVM model performs slightly better than the Logistic Regression model on both accuracy and macro-F1, but the difference is not very large (e.g., 0.01 in accuracy and macro-F1).

Comparing the two models in Table 3, we can see that the Linear SVM model has a slightly higher accuracy and macro-F1 score compared to the Logistic Regression model on the test set. However, both models perform quite well, with accuracies and macro-F1 scores above 0.90, indicating that the TF-IDF features are effective for this classification task.

	Precision	Recall	F1-Score	Support
1	0.94	0.90	0.92	1900
2	0.95	0.98	0.97	1900
3	0.89	0.88	0.89	1900
4	0.89	0.90	0.90	1900
Accuracy			0.92	7600
Macro avg	0.92	0.92	0.92	7600
Weighted avg	0.92	0.92	0.92	7600

Table 1: Classification report for Linear SVM on test set where classes are labeled as 1: world, 2: sports, 3: business, and 4: sci/tech.

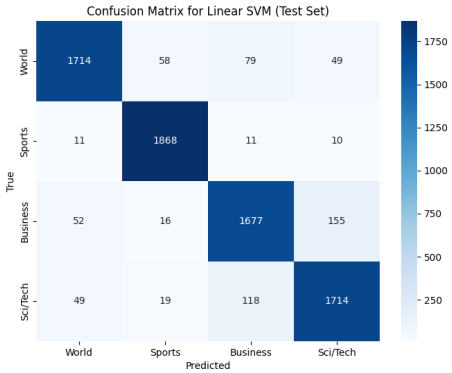

Figure 2: Confusion Matrix for Linear SVM on Test Set

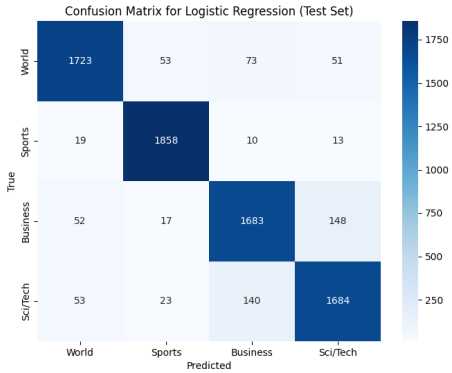

Figure 3: Confusion Matrix for Logistic Regression on Test Set

	Precision	Recall	F1-Score	Support
1	0.93	0.91	0.92	1900
2	0.95	0.98	0.96	1900
3	0.88	0.89	0.88	1900
4	0.89	0.89	0.89	1900
Accuracy			0.91	7600
Macro avg	0.91	0.91	0.91	7600
Weighted avg	0.91	0.91	0.91	7600

Table 2: Classification report for Logistic Regression on test set where classes are labeled as 1: world, 2: sports, 3: business, and 4: sci/tech.

Model	Accuracy	Macro-F1
Linear SVM	**0.92**	**0.92**
Logistic Regression	0.91	0.91

Table 3: Comparison of Accuracy and Macro-F1 for Linear SVM and Logistic Regression on Test Set

5 Error Analysis

In this section, we will analyze the errors made by our best-performing models (Linear SVM and Logistic Regression) on the test set. We will look at the confusion matrices and t-SNE projections of the predictions to understand which classes are most often confused and why.

Looking at Figures 2 and 3, we can see that the TF-IDf features form many distinct clusters in the t-SNE projection, the classes seem to be mostly seperable, but there is some overlap in the center of the projection. The confusion matrices show that the most common misclassification is between the sci/tech and business classes, which is reflected in the t-SNE projection where the clusters for these classes overlap more than the others. This suggests that there are some samples in the test set that have features that are similar to both classes, which might explain why they are misclassified.

To further analyze this, we recorded 20 misclassified samples for each of the two most common confusions (i.e., business predicted as sci/tech, and sci/tech predicted as business) and looked at the raw text to see if there are any common features that might explain the confusion. We found that the errors can be put into 5 categories: 1) labeling errors, 2) ambiguous samples, 3) samples with overlapping keywords, 4) samples with insufficient context, and 5) samples that are simply difficult to classify.

Labeling errors account for a small but notable portion of the misclassified samples. These are articles that appear to have been assigned the wrong category in the AG News dataset itself, independent of any model behavior. For example, the article "Teenage T. rex's monster growth" describes paleontological research and is clearly a science article, yet it carries a Business label in the dataset. Similarly, the Hurricane Charley article concerns meteorological forecasting accuracy and has no apparent business relevance. Since AG News was constructed through automated web scraping and keyword-based labeling rather than human annotation, a degree of label noise is expected. These errors represent a hard ceiling on achievable accuracy that no model improvement can overcome, as the ground truth itself is incorrect.

Ambiguous samples are articles that could reasonably belong to either category, where even a human annotator might disagree on the correct label. The Amazon/Joyo acquisition is a representative example: it is structurally an M&A story, which would typically fall under Business, yet it concerns an internet retail platform and involves companies whose primary identity is technological. Similarly, the Yahoo pricing article discusses a commercial strategy decision but centers on web hosting and domain registration products. These cases reflect a fundamental limitation of the AG News label schema, which forces a single category onto articles that genuinely straddle two domains. A model trained on this schema will inevitably produce errors on such samples regardless of its representational capacity.

Samples with overlapping keywords are perhaps the most common error type and the most directly attributable to the TF-IDF representation. These are articles whose true category is clear, but whose vocabulary is dominated by terms strongly associated with the other category. Tech company earnings reports are the clearest example: articles about Intuit, Ciena, and Microsoft contain financial language such as "quarterly loss", "forecast", and "government deal" that the model associates with Business, while the entities themselves (software companies, telecom equipment makers) carry strong Sci/Tech signal. The model has no mechanism to weight the framing of an article over the entities it mentions, so whichever signal is stronger in the term frequency wins.

Samples with insufficient context are short or sparsely worded articles where the preprocessed text does not contain enough discriminative signal to make a confident prediction. The Nortel joke column is an extreme case — after preprocessing, almost no meaningful content remains, and the only signal is a stock ticker. The SUV safety article is similarly sparse after stopword removal, leaving terms like "traffic crash fatality data" and "sport utility vehicle" that do not map cleanly to either Business or Sci/Tech. These errors would likely persist even with a contextual model, since the input itself does not contain sufficient information to determine the correct category.

Finally, some samples are simply difficult to classify and do not fit neatly into any of the above categories. The S2io/Sun Microsystems article, for instance, describes a technology partnership in terms that are equally weighted toward business development and technical infrastructure. The HP holiday shoppers article frames a product strategy announcement in consumer marketing language while discussing flat-screen TVs and media devices. These articles are well-written, correctly labeled, and contain sufficient context, yet the features available to a bag-of-words model do not cleanly distinguish them from the opposite category. Addressing these cases would likely require a model capable of understanding discourse structure and authorial intent, such as a fine-tuned transformer, rather than improvements to the feature representation.

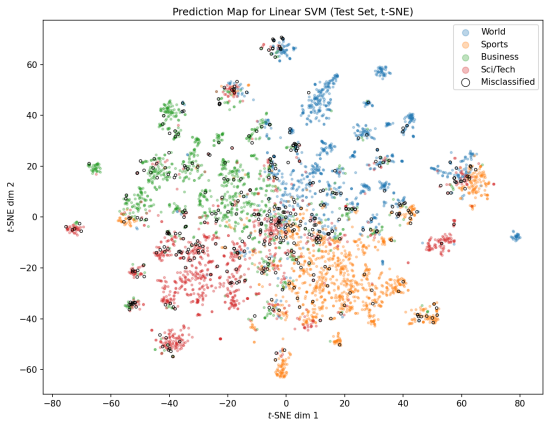

Figure 4: t-SNE Projection of Linear SVM Predictions on Test Set. Correctly classified samples are colored according to their true class, while misclassified samples are outlined in black.

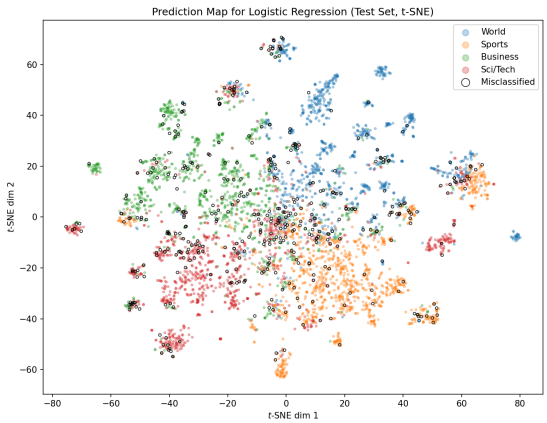

Figure 5: t-SNE Projection of Logistic Regression Predictions on Test Set. Correctly classified samples are colored according to their true class, while misclassified samples are outlined in black.

6 Discussion

7 Limitations

References

[Zhang et al.(2015)Zhang, Zhao, and LeCun] Xiang Zhang, Junbo Zhao, and Yann LeCun. Character-level convolutional networks for text classification. In *Advances in Neural Information Processing Systems 28 (NIPS 2015)*, 2015.

A Appendix

Table 4: Misclassified samples: *Business* articles predicted as *Sci/Tech*.

Preprocessed Text	**Raw Text (Excerpt)**
teenage rex monster growth tyrannosaurus rex achieved massive size due enormous growth spurt adolescent year	Teenage T. rex's monster growth: Tyrannosaurus rex achieved its massive size due to an enormous growth spurt during its adolescent years.
people eligible get google ipo google billed ipo way everyday people...	Some People Not Eligible to Get in on Google IPO: Google has billed its IPO as a way for everyday people to get in on the process...
rival try turn table charles schwab michael liedtke san francisco ap low price iconoclastic attitude discount stock broker...	Rivals Try to Turn Tables on Charles Schwab: With its low prices and iconoclastic attitude, discount stock broker Charles Schwab Corp. represented an annoying stone in Wall Street's wing-tipped shoes...
dell exit lowend china consumer pc market hong kong reuters dell inc world largest pc maker said monday left lowend consumer pc market china...	Dell Exits Low-End China Consumer PC Market: Dell Inc., the world's largest PC maker, said on Monday it has left the low-end consumer PC market in China...
ibm buy two danish service firm ibm said tuesday acquired pair danish service firm part effort broaden presence scandinavia...	IBM Buys Two Danish Services Firms: IBM said Tuesday it has acquired a pair of Danish IT services firms as part of its effort to broaden its presence in Scandinavia...
microsoft finalises threeyear government deal hot heel 10year strategic partnership london borough newham...	Microsoft finalises three-year government deal: Hot on the heels of its 10-year strategic partnership with the London Borough of Newham...
intuit post wider loss charge reuters intuit inc maker ofthe 1 u tax presentation software turbotax wednesdayposted wider quarterly loss...	Intuit Posts Wider Loss After Charge: Intuit Inc., maker of the No. 1 U.S. tax presentation software TurboTax, posted a wider quarterly loss...
hurricane center 39s projection charley far data show fort lauderdale fla krt despite criticism better anticipated hurricane charley...	Hurricane center's projection on Charley not far off, data show: Despite criticism that it should have better anticipated Hurricane Charley's rapid intensification...
ciena post loss forecast flat sale chicago reuters telecommunication equipment maker ciena corp reported wider loss thefiscal third quarter...	Ciena Posts a Loss, Forecasts Flat Sales: Telecommunications equipment maker Ciena Corp. reported a wider loss for the fiscal third quarter due to slack demand...
report consumer tuning plasma tv firstquarter shipment plasma television united state doubled previous year according research firm isuppli...	Report: Consumers tuning in to plasma TVs: First-quarter shipments of plasma televisions in the United States more than doubled from the previous year...

Table 5: Misclassified samples: *Sci/Tech* articles predicted as *Business*.

Preprocessed Text	**Raw Text (Excerpt)**
intel delay product aimed highdefinition tv san francisco latest series product delay intel corp postponed launch video display chip...	Intel to delay product aimed for high-definition TVs: In the latest of a series of product delays, Intel Corp. has postponed the launch of a video display chip...
yahoo ups ante small business web giant yahoo gambling price cut domain name registration web hosting product make competitive discounter...	Yahoo! Ups Ante for Small Businesses: Web giant Yahoo! is gambling that price cuts on domain name registration and Web hosting products will make it more competitive...
suv live let die new york newly released traffic crash fatality data something everyone debate safety sport utility vehicle	SUVs: Live And Let Die: The newly released traffic crash fatality data have something for everyone in the debate about the safety of sport utility vehicles.
ross store profit plummet 40 percent ap discount retailer ross store inc wednesday said profit fell 40 percent latest quarter due problem new computer system...	Ross Stores Profit Plummets 40 Percent: Discount retailer Ross Stores Inc. said its profit fell about 40 percent in the latest quarter due to problems with a new computer system...
ohio sue best buy alleging used sale ap ohio authority sued best buy co inc thursday alleging electronics retailer engaged unfair deceptive business practice	Ohio Sues Best Buy, Alleging Used Sales: Ohio authorities sued Best Buy Co. Inc. alleging the electronics retailer engaged in unfair and deceptive business practices.
amazon snap china largest web retailer newsfactor amazoncom nasdaq amzn said buy joyocom limited british virgin island company operates largest internet retail web site china...	Amazon Snaps Up China's Largest Web Retailer: Amazon.com has said it will buy Joyo.com Limited, a company that operates the largest Internet retail Web site in China, for $75 million.
deal s2io champing gigabit ottawa local firm say help shrink backup time large data centre growing business thanks alliance sun microsystems inc	Deal has S2io champing at the gigabit: A local firm that says it can help shrink backup times at large data centres is growing its business thanks to an alliance with Sun Microsystems Inc.
mom 2005 released manufacturing microsoft wednesday announced release manufacturing microsoft operation manager mom 2005 workgroup edition...	MOM 2005 Released to Manufacturing: Microsoft announced the release to manufacturing of Microsoft Operations Manager (MOM) 2005 and MOM 2005 Workgroup Edition...
hp tempt holiday shopper sight sound computerhardware giant best known product pc printer laid plan become brandname consumer electronics product...	HP to tempt holiday shoppers with sights and sounds: The computer-hardware giant laid out its plan to become a brand-name in consumer electronics products such as flat-screen TVs and music players...
canadian invasion week ago story nortel nyse nt asked people submit canadian joke good place reveal winner	A Canadian Invasion: A few weeks ago, in a story on Nortel (NYSE: NT), I asked people to submit a Canadian joke to me.